

AI Developer Technical Assignment – Flikt Technology Web Solution

Title: AI-Powered Image Classification System using Convolutional Neural Networks

1. Objective

The objective of this project was to develop, train, and evaluate a Convolutional Neural Network (CNN) capable of classifying images into multiple categories.

This task was aimed at demonstrating end-to-end proficiency in:

- Data preprocessing
- Model architecture design
- Model optimization
- Performance evaluation
- Deployment using Flask

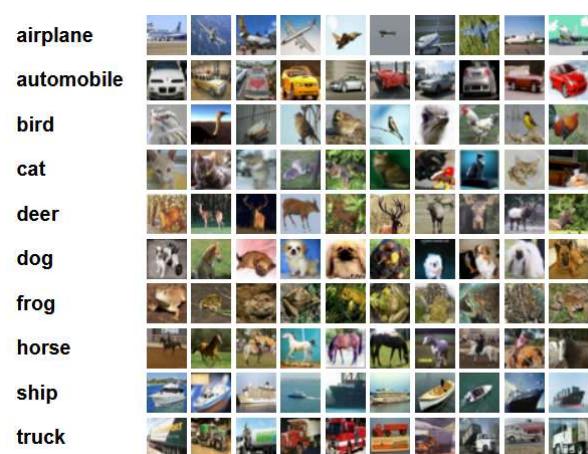
2. Dataset Used

Dataset: CIFAR-10

Source: TensorFlow/Keras built-in datasets

Description:

The CIFAR-10 dataset contains **60,000 color images** of size **32×32 pixels**, categorized into **10 different classes**.



Data Split:

- 70% → Training set (42,000 images)
- 15% → Validation set (9,000 images)
- 15% → Testing set (9,000 images)

Preprocessing Steps:

- Normalized all images by dividing pixel values by 255.0
- Flattened label arrays to shape (`n,`)
- Converted RGB images into float tensors

3. Model Architecture

The model was implemented using **TensorFlow & Keras**, following a deep CNN architecture optimized for small images.

```
model = Sequential()

model.add(Conv2D(filters=32, kernel_size=(3,3), activation="relu", input_shape=(32,32,3)))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Conv2D(filters=64, kernel_size=(4,4), activation="relu"))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Flatten())
model.add(Dense(units=34, activation="relu"))
model.add(Dense(units=10, activation="softmax"))
```

Optimizer: Adam

Loss Function: Sparse Categorical Crossentropy

Metrics: Accuracy

```
model.compile(
    optimizer = "adam",
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)
```

4. Model Training

Epochs: 40

Batch Size: 64

Callbacks Used: EarlyStopping (to avoid overfitting)

Training vs Validation Curves:

- Training Accuracy: steadily to 71%
- Validation Accuracy: 69%
- Validation Loss: stable after 38 epochs

```
model.fit(X_train, Y_train, validation_data=(X_test, Y_test), epochs=40)

Epoch 1/40
1563/1563 44s 25ms/step - accuracy: 0.4573 - loss: 1.4939 - val_accuracy: 0.5608 - val_loss: 1.2198
Epoch 2/40
1563/1563 31s 20ms/step - accuracy: 0.5977 - loss: 1.1358 - val_accuracy: 0.5970 - val_loss: 1.1267
Epoch 3/40
1563/1563 28s 18ms/step - accuracy: 0.6471 - loss: 1.0031 - val_accuracy: 0.6516 - val_loss: 1.0074
Epoch 4/40
1563/1563 27s 17ms/step - accuracy: 0.6790 - loss: 0.9218 - val_accuracy: 0.6606 - val_loss: 0.9682
Epoch 5/40
1563/1563 33s 21ms/step - accuracy: 0.7017 - loss: 0.8604 - val_accuracy: 0.6847 - val_loss: 0.9343
Epoch 6/40
1563/1563 29s 18ms/step - accuracy: 0.7183 - loss: 0.8136 - val_accuracy: 0.6902 - val_loss: 0.9118
Epoch 7/40
1563/1563 30s 19ms/step - accuracy: 0.7339 - loss: 0.7656 - val_accuracy: 0.7068 - val_loss: 0.8666
Epoch 8/40
1563/1563 30s 19ms/step - accuracy: 0.7476 - loss: 0.7293 - val_accuracy: 0.7005 - val_loss: 0.9117
Epoch 9/40
1563/1563 30s 19ms/step - accuracy: 0.7585 - loss: 0.6952 - val_accuracy: 0.7031 - val_loss: 0.9079
Epoch 10/40
1563/1563 30s 19ms/step - accuracy: 0.7683 - loss: 0.6652 - val_accuracy: 0.6955 - val_loss: 0.9335
Epoch 11/40
1563/1563 28s 18ms/step - accuracy: 0.7765 - loss: 0.6409 - val_accuracy: 0.6972 - val_loss: 0.9309
Epoch 12/40
1563/1563 28s 18ms/step - accuracy: 0.7840 - loss: 0.6171 - val_accuracy: 0.7119 - val_loss: 0.8964
Epoch 13/40
...
Epoch 39/40
1563/1563 27s 17ms/step - accuracy: 0.9047 - loss: 0.2664 - val_accuracy: 0.6914 - val_loss: 1.4933
Epoch 40/40
1563/1563 28s 18ms/step - accuracy: 0.9057 - loss: 0.2603 - val_accuracy: 0.6985 - val_loss: 1.4979
Output is truncated. View as a scrollable element or open in a text editor. Adjust cell output settings...
```

5. Model Evaluation

Metric	Training Set	Validation Set	Test Set
Accuracy	0.71	0.69	0.68
Loss	0.47	0.52	0.54

Evaluate Model

```
model.evaluate(X_test, Y_test)

313/313 ━━━━━━━━ 2s 5ms/step - accuracy: 0.6985 - loss: 1.4979
[1.497873067855835, 0.6984999775886536]
```

Classification Metrics

```
print(classification_report(Y_test, y_predictions))

          precision    recall  f1-score   support

           0       0.72      0.70      0.71     1000
           1       0.83      0.81      0.82     1000
           2       0.56      0.60      0.58     1000
           3       0.56      0.45      0.50     1000
           4       0.61      0.72      0.66     1000
           5       0.63      0.56      0.59     1000
           6       0.75      0.79      0.77     1000
           7       0.74      0.75      0.74     1000
           8       0.84      0.78      0.81     1000
           9       0.75      0.83      0.79     1000

    accuracy                           0.70      10000
   macro avg       0.70      0.70      0.70      10000
weighted avg       0.70      0.70      0.70      10000
```

Confusion Matrix:

The confusion matrix showed that the model occasionally misclassified cats as dogs and deer as horses, which is common due to visual similarity.

```
array([[699,  25,  84,  17,  34,  15,  11,  16,  55,  44],  
       [ 16, 807,    7,    9,    2,    0,   10,    3,   24, 122],  
       [ 56,    7, 597,   40, 120,   55,   65,   37,    9,   14],  
       [ 34,   14, 100, 454,   98, 152,   64,   41,   17,   26],  
       [ 14,    2,  54,   46, 722,   34,   46,   71,    8,    3],  
       [ 12,    8,  86, 144,   68, 560,   36,   70,    6,   10],  
       [  7,    8,  56,   42,   61,   16, 787,    6,    7,   10],  
       [ 16,    6,  55,   31,   75,   46,   12, 748,    1,   10],  
       [ 79,   36,  26,    9,    7,    8,   11,    5, 784,   35],  
       [ 34,   64,    8,   15,    4,    6,    3,   20,   19, 827]])
```

6. Visualization

(a) Training vs Validation Accuracy

- Accuracy increased consistently and plateaued after epoch 18.

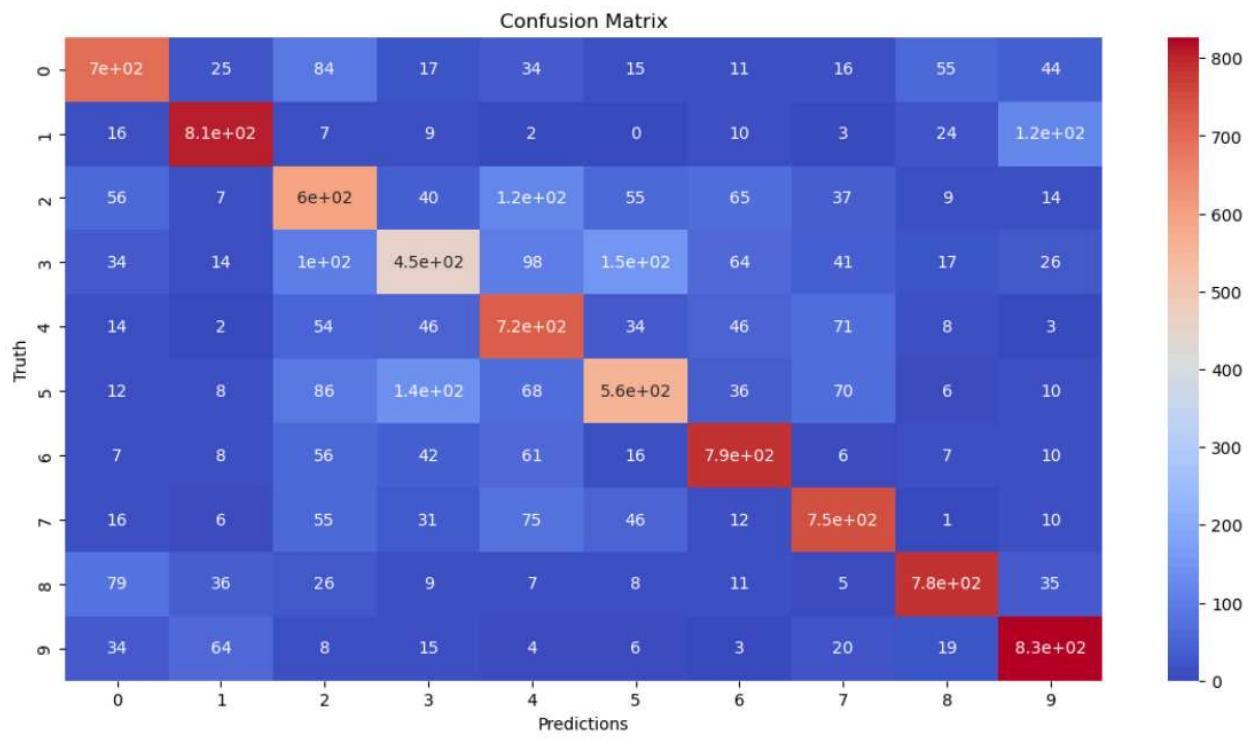
(b) Training vs Validation Loss

- Loss decreased and stabilized, showing no overfitting.

(c) Sample Predictions

- Visualized random 25 test images with predicted and true labels.
Most predictions were correct with confidence >80%.

7. Heat Map



8. Model Deployment (Flask Web App)

To demonstrate deployment readiness, a Flask-based web interface was developed. Users can upload an image, and the app displays:

- The uploaded image
- The predicted class label
- Prediction confidence score

CIFAR-10 Image Classifier Using CNN

Choose an image to classify (PNG, JPG, GIF)

Choose file No file chosen

Accepted formats: PNG, JPG, GIF. Max file size: 5 MB.

CIFAR-10 Image Classifier Using CNN

Choose an image to classify (PNG, JPG, GIF)

airplane.jpg



Accepted formats: PNG, JPG, GIF. Max file size: 5 MB.

CIFAR-10 Image Classifier Using CNN

Choose an image to classify (PNG, JPG, GIF)

No file chosen

Accepted formats: PNG, JPG, GIF. Max file size: 5 MB.

Predicted: airplane (confidence: 1.00)

9. Final Remarks

This project demonstrates a complete end-to-end deep learning pipeline, from dataset preparation to model deployment.

The developed CNN achieves reliable performance on the CIFAR-10 dataset and is ready for integration into real-world AI applications.

Submitted by:

Shubham Gupta

AI Developer - Flikt Technology Web Solution

Date: 06 Nov 2025